PROJECT STATUS REPORT

Claude Code Advancements

Research, tools, and systems for AI-assisted development

10,706	1,902,376	9	1381
TESTS PASSING	LINES OF CODE	MODULES	COMMITS

51	461	18	41d
master tasks	intelligence findings	live hooks	project age

Session 246 · 2026-04-01

github.com/mpshields96/ClaudeCodeAdvancements

CONTENTS

Contents

1 Executive Summary 3
Project overview, health dashboard, and key metrics

2 Five Frontiers 4
Status of the five core research areas

3 Module Deep-Dives 5
Detailed breakdown of each module's capabilities and progress

4 Master Tasks 9
All aspirational goals — complete, active, pending, and blocked

5 Priority Queue 15
Decay-based priority scoring for active work

6 Live Infrastructure 16
Hook architecture integrated into Claude Code

7 Intelligence & Research 16
Reddit findings, academic papers, and self-learning metrics

8 Architecture Decisions 22
Key technical decisions and their rationale

9 Risks & Blockers 23
Known issues, mitigations, and technical debt

10 Next Priorities 23
Upcoming work items ranked by impact

11 Honest Assessment 24
Objective gaps, limitations, and areas falling short

EXECUTIVE SUMMARY

Executive Summary

ClaudeCodeAdvancements is a research and development project building validated tools, hooks, and systems for Claude Code users and AI-assisted development workflows. Launched 2026-02-19 with 246 development sessions completed. All five original research frontiers are production-complete with 10,706 automated tests across 9 modules. The project has expanded into 51 master-level tasks covering autonomous intelligence gathering, self-learning loops, academic research integration, and professional design output. Zero external Python dependencies.

SINCE LAST REPORT — Session 169 → 246

+1205	+1732063	+391	+1	+49	+1
Tests	LOC	Commits	Completed MTs	Findings	Delivered

Newly Completed: MT-37 — AI-Driven Investment Research & Portfolio Intelligence

Kalshi P&L: +$32.99 (WR: +1.5pp)

Project Health

TESTS

10,706 / 10,706

282 suites — 100% pass rate

CODEBASE

1,902,376

Source: 1,479,758 Test: 422,618

Test-to-source ratio: 0.29:1

VELOCITY

1381 commits

33.7/day · 5.6/session

41 days · 246 sessions

MODULES

9

4643 source files · 904 test files

0 stubs · 0 syntax errors

MASTER TASKS

51

10 done · 35 active · 5 pending

INTELLIGENCE

461 findings

1242 papers · 18 live hooks

Zero external dependencies

FIVE FRONTIERS

Five Frontiers

The five core research areas that drive CCA's development, each validated through community intelligence, Anthropic research, and developer surveys.

1 Persistent Cross-Session Memory

Active CRITICAL

Every session starts from zero — this gives Claude a brain that persists.

TESTS: **340** LOC: **1,909**

2 Spec-Driven Development

Active HIGH

Unstructured prompting produces poor architecture. Enforce requirements > design > tasks > implement.

TESTS: **205** LOC: **1,191**

3 Context Health Monitor

Active HIGH

Context rot causes silent output degradation. Monitor, warn, and auto-handoff.

TESTS: **434** LOC: **3,383**

4 Multi-Agent Conflict Guard

Active HIGH

Parallel agents overwrite each other. Guard credentials, paths, and system modifications.

TESTS: **1102** LOC: **6,964**

5 Usage Transparency Dashboard

Active MEDIUM

No real-time token/cost visibility. Counter, alerts, and structural completeness.

TESTS: **369** LOC: **3,030**

MODULE DEEP-DIVES

Module Deep-Dives

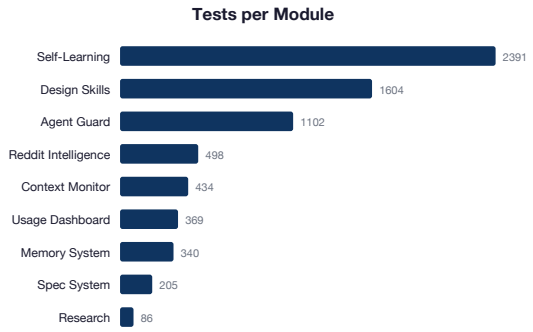

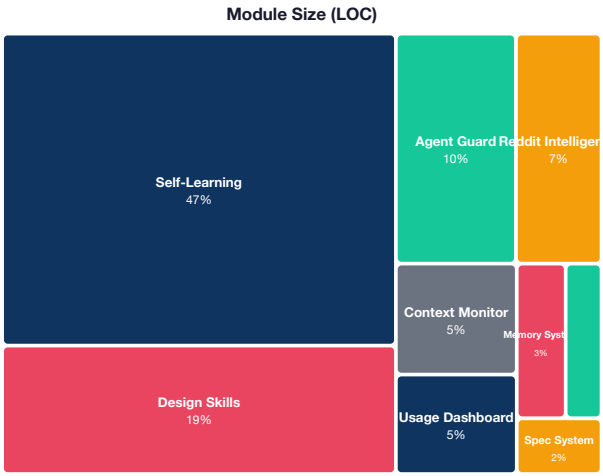

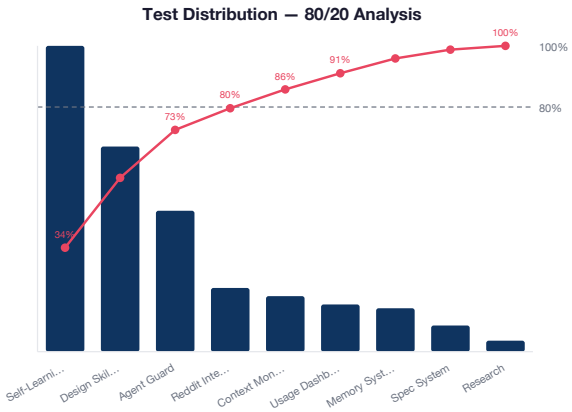

Memory System

Active

TESTS **340** LOC **1,909** FILES **5** `memory-system/`

Persistent cross-session memory. Captures decisions, preferences, and context across Claude Code sessions using hooks and MCP.

KEY COMPONENTS

- Capture hook (PostToolUse + Stop)
- CLI viewer (stats, search, list)
- 8-char hex ID schema
- MCP retrieval server
- Compaction-resistant handoff

WHY THIS MATTERS Means Claude remembers your preferences and past decisions across sessions — no more re-explaining context every time you start a new chat.

Spec System

Active

TESTS **205** LOC **1,191** FILES **4** `spec-system/`

Spec-driven development workflow. Enforces requirements > design > tasks > implement sequence with approval gates.

KEY COMPONENTS

- /spec:requirements interview scaffold
- /spec:tasks atomic decomposer
- PreToolUse spec validator hook
- /spec:design architecture generator
- /spec:design-review (4 expert personas)
- UserPromptSubmit auto-activation

WHY THIS MATTERS Forces Claude to plan before coding. You get a reviewable blueprint (requirements, design, tasks) instead of spaghetti code from vague prompts.

Context Monitor

Active

TESTS **434** LOC **3,383** FILES **11** `context-monitor/`

Context health monitoring with adaptive thresholds. Tracks token usage, warns at quality ceilings, blocks exit at critical levels.

KEY COMPONENTS

- PostToolUse token meter
- Stop hook auto-handoff
- Automatic session wrap trigger
- PreToolUse threshold alert
- Compact anchor (every 10 turns)
- Session pacer (2-3h autonomous runs)

WHY THIS MATTERS Warns you when Claude's memory is getting full so it doesn't silently start forgetting instructions mid-session. Prevents the 'why did it ignore my rules?' problem.

Agent Guard

Active

TESTS **1102** LOC **6,964** FILES **24** `agent-guard/`

Multi-agent safety, conflict prevention, and Senior Dev code review. Guards credentials, network ports, dangerous paths, and provides automated quality scoring on every write.

KEY COMPONENTS

- Mobile approver (ntfy.sh push)
- Network/port exposure guard
- Path validator (LIVE in hooks)
- File ownership manifest
- Bash command safety guard (LIVE in hooks)
- PR effort scorer (1-5 scale)
- False positive filter
- Tech debt tracker (trend analysis)
- Credential extraction guard
- Content scanner (9 threat categories)
- Session guard (slop detection)
- Edit retry guard (LIVE in hooks)
- SATD detector (TODO/FIXME/HACK markers)
- Code quality scorer (0-100, A-F)
- Review classifier (6 categories)
- ADR reader (architecture decision discovery)

- Senior Dev orchestrator hook (LIVE)

WHY THIS MATTERS Stops Claude from accidentally deleting files, exposing API keys, or breaking your computer when running autonomously overnight.

Usage Dashboard

Active

TESTS **369** LOC **3,030** FILES **7** `usage-dashboard/`

Token and cost transparency. CLI counter, OpenTelemetry receiver, cost alerts, and structural completeness checker.

KEY COMPONENTS

- Token counter CLI
- PreToolUse cost alert hook
- OTLP HTTP/JSON receiver
- /arewedone completeness checker

WHY THIS MATTERS Shows you exactly how many tokens and dollars each session costs, so you can spot waste and stay within your Max subscription limits.

Reddit Intelligence

Complete

TESTS **498** LOC **4,964** FILES **9** `reddit-intelligence/`

Community signal research. Discovers tools, patterns, and pain points from Reddit and GitHub to drive development priorities.

KEY COMPONENTS

- Reddit reader (posts + all comments)
- GitHub repo evaluator
- Subreddit profile registry (10 profiles)
- Autonomous scan pipeline
- Sandboxed repo tester

WHY THIS MATTERS Automatically finds the best new tools, tips, and techniques from Reddit so you don't have to scroll through hundreds of posts yourself.

Self-Learning

Complete

TESTS **2391** LOC **30,981** FILES **85** `self-learning/`

Cross-session improvement through structured observation, pattern detection, autonomous proposals, and research ROI tracking.

KEY COMPONENTS

- JSONL event journal
- Trace analyzer (4 detectors)
- Sentinel adaptive mutation
- Findings resurfacer
- Overnight detector (time-stratified analysis)
- Trade reflector (5 pattern detectors)
- Pattern reflector + strategy engine
- YoYo improvement loop
- Academic paper scanner
- Skillbook injection hook (LIVE)
- Research outcomes ROI tracker

WHY THIS MATTERS Makes your Kalshi bot smarter over time by analyzing what worked and what didn't. The bot learns from its own wins and losses instead of repeating mistakes.

Design Skills

Active

TESTS **1604** LOC **12,765** FILES **19** `design-skills/`

Professional visual output. Typst-based PDF reports, slides, HTML dashboards, SVG charts, website generation, and daily project snapshots.

KEY COMPONENTS

- Report generator (Typst pipeline)
- Dashboard generator (HTML)
- Website generator (landing + docs)
- Design guide + visual language
- Slide generator (16:9 PDF)
- Chart generator (SVG)
- Daily snapshot engine (metric diffs)

WHY THIS MATTERS Generates professional PDF reports and charts automatically — no manual formatting. One command gives you a polished status overview you could show anyone.

Research

Complete

TESTS **86** LOC **1,397** FILES **3** `research/`

R&D tools and experimental capabilities including iOS project generation and Xcode build wrapper.

KEY COMPONENTS

- iOS project generator (SwiftUI)
- Xcode build wrapper

WHY THIS MATTERS R&D sandbox for experimental features. Things get prototyped here before becoming real modules — keeps the main codebase clean.

MASTER TASKS

Master Tasks

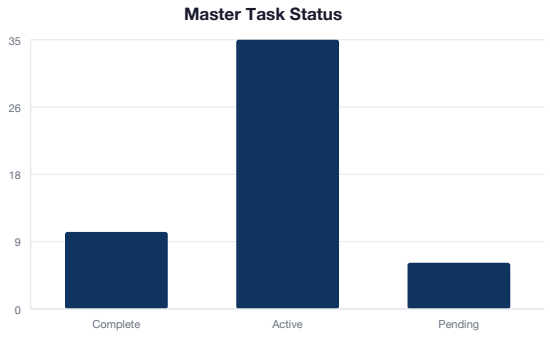

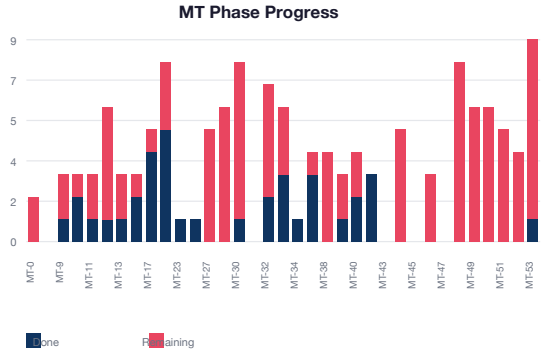

51 tasks defined · **10 complete** · **35 active** · 5 pending · 1 blocked

Completed

10 tasks delivered

MT–2 Mermaid Architecture Diagrams in Spec System

Complete

Auto-generates architecture diagrams in the spec workflow. See the design visually before writing code.

MT–3 Virtual Design Team Plugin

Complete

Multiple AI personas (backend, frontend, security) review designs from different expert perspectives.

MT–4 Frontend Design Excellence via Design Vocabulary

Complete

Gives Claude a professional design vocabulary so UI recommendations use real design principles, not vague suggestions.

MT–6 On-Demand Subreddit Scanner

✓ 43 tests Complete

Scan any subreddit on demand for Claude Code relevant posts. The 'nuclear' option for deep intelligence sweeps.

profiles.py: 10 builtin subreddit profiles (4 domains: claude/trading/dev/res... · --profile and --quick flags added to nuclear_fetcher.py

MT–7 Programmatic Trace Analysis for Self-Learning

✓ 50 tests Complete

Finds patterns in how Claude wastes tokens (retries, unnecessary reads) so sessions run faster and cheaper.

self-learning/trace_analyzer.py: TranscriptEntry, TranscriptSession, RetryDet...

MT-20 **Senior Developer Agent** ✓ 890 tests Complete

An automated code reviewer that catches bugs, tech debt, and design issues before they become problems. Like having a senior engineer on call 24/7.

MT-22 **Autonomous 1-Hour Loop** ✓ 66 tests Complete

Makes the whole system run without you. Sessions start, do work, wrap up, and hand off to the next session automatically.

MT-26 **Financial Intelligence Engine** ✓ 43 tests Complete

Financial intelligence tools for the Kalshi bot: regime detection, calibration analysis, order flow. Directly improves bet quality.

MT-29 **Cowork + Claude Pro Bridge Hivemind** Complete

Bridge between Claude Pro web Cowork mode and Claude Code CLI for hivemind coordination across interfaces.

MT-37 **AI-Driven Investment Research & Portfolio Intelligence** ✓ 250 tests Complete

Long-term investment research engine — ETF analysis, portfolio construction, factor models. Same rigor as Kalshi: academic basis + backtesting + math validation.

Phase 1 COMPLETE (S178-180): MT37_RESEARCH.md — 1762 lines, 42 papers synthes...

Active

35 tasks in progress

MT-0 **Kalshi Bot Self-Learning Integration** ✓ 24 0/2

COMPLETE (Phase 1: S21, Phase 2: deployed via polymarket-bot monitoring chat).

MT-1 **Maestro-Style Visual Grid UI** Active

MOSTLY SELF-RESOLVED (S96). Multiple tools now exist:

MT-9 **Autonomous Cross-Subreddit Intelligence Gathering** ✓ 54 1/3

Phase 2 COMPLETE (Session 31). Delivered:

MT-10 **YoYo-Pattern Continuous Self-Learning and Self-Building** ✓ 64 2/3

Phase 1 COMPLETE (Session 28), QualityGate added (Session 29), E2E validated (Session 29).

MT–11 Autonomous GitHub Repository Intelligence

✓ 62 1/3

Phase 2 COMPLETE (Session 83). Delivered:

MT–12 Academic Research Paper Integration

✓ 89 1/6

Phase 4 COMPLETE (Session 101). Delivered:

MT–13 iOS App Development Capability

✓ 57 1/3

Phase 2 COMPLETE (Session 49). Delivered:

MT–14 Autonomous Re-Scanning of Previously Scanned Subreddits

2/3

Phase 1 COMPLETE (Session 35). Delivered:

MT–17 UI/Design Excellence and Professional Report Generation

✓ 42 4/5

Phase 1 COMPLETE (Session 41). Delivered:

MT–21 Hivemind Multi-Chat Coordination

✓ 69 5/8

Infrastructure COMPLETE. Phase 1 validation IN PROGRESS (S86). **→ 1-3 more validation tests, then Phase 1 gate assessment.**

MT–23 Mobile Remote Control v2

Active

EXTERNALLY RESOLVED (S112). Claude Code Channels shipped March 20, 2026. **→ Install Telegram channel plugin and test with Kalshi bot session.**

MT–24 Visualization & Graphics Engine

✓ 41 1/1

Phase 1 COMPLETE (S111) — `trading_chart.py` (5 SVG chart types: PnL curve, win rate, strategy matrix, drawdown, heatmap). 41 tests.

MT–27 CCA Nuclear v2

✓ 27 0/5

Phases 1-5 COMPLETE.

MT–28 Self-Learning v2

✓ 73 0/6

COMPLETED (S111). All 6 phases done:

MT-30 Session Daemon

✓ 116 1/8

Phase 8 COMPLETE (S128). S129: +preflight command, +rich --status with audit log, +AUTOLOOP_SETUP.md. Ready for live supervised dry run.

MT-31 Gemini Pro Integration

Active

S125 E2E validation: Gemini Flash WORKS via @rlabs-inc/gemini-mcp (v0.8.1). Pro UNAVAILABLE — Google One gives web app Pro access only, not API access. API Pro requires separate Google Cloud billing. Matthew directive: Flash only (Option 1). Scope narrowed to Gemini Flash integration.

MT-32 Visual Excellence & Design Engineering

✓ 890 2/7

Phase 3 COMPLETE (S231). Phases 1-3 done. Next: Phase 4 (UI component library).

MT-33 Strategic Intelligence Report

3/6

Unknown

MT-34 Medical AI Tool — OpenEvidence Replacement

Active

IDEA LOGGED — Matthew needs more time to refine the concept. Do NOT start building until Matthew gives the green light.

MT-35 Background Autoloop — Non-Intrusive Desktop Loop

✓ 23 3/4

- Phase 1 COMPLETE (S142): save/restore frontmost app around autoloop trigger (~3-5s takeover)

MT-38 Peak/Off-Peak Token Budget System

✓ 21 0/4

Phase 1-3 COMPLETE (S154-S155). Rule active globally. token_budget.py built (21 tests, S154). Peak-hour agent blocking hook active (S155). Phase 4 (autoloop scheduling) remaining.

MT-39 Priority Picker Overhaul — Dust Detection

1/3

Phase 1 COMPLETE (S160). 93 existing tests still pass. `python3 priority_picker.py dust` now shows 15 dusty MTs.

MT-40 Automated Nuclear Scanning Loop

✓ 17 2/4

Phase 1-2 COMPLETE (S160). Phase 3-4 remaining.

MT-41 Synthetic MT Origination

✓ 53 3/3

COMPLETE (S160-S163). All 3 phases delivered. 53 tests total.

MT–43 Data Analysis Excellence

Active

IDEA LOGGED (S169). Pending definition of specific quality gaps vs current MCP tool output.

MT–44 Automated Cross-Chat Task Execution Pipeline

0/5

IDEA LOGGED (S169). Infrastructure partially built. Gaps documented in ORCHESTRATION_GAPS.md.

MT–45 USAGE-5 Streamlit/Web Dashboard

Active

IDEA LOGGED (S169). Low priority — current CLI and HTML dashboards cover most needs. Build only if Matthew specifically requests real-time web visualization.

MT–46 Report Self-Audit System

0/3

IDEA LOGGED (S169). doc_drift_checker.py and arewedone.py provide the foundation.

MT–47 External Tool Evaluation Pipeline

Active

IDEA LOGGED (S169). 7+ tools in backlog. Safety-first: read source only, rebuild from scratch if useful.

MT–48 Report Visual Polish

0/8

IDEA LOGGED (S169). All issues are concrete and actionable — no research needed, just implementation.

MT–49 Self-Learning Evolution Engine

0/6

UBER-LEVEL DOCUMENTED (S169). Permanent MT — no completion date. Existing infrastructure is strong. Focus on meta-learning and automated principle discovery as first priorities.

MT–50 Kalshi Copytrade Bot Research

0/6

RESEARCH PHASE. S181b directive documented. No implementation yet.

MT–51 Kalshi New Market Expansion Research

0/5

RESEARCH PHASE. S181b directive documented. No implementation yet.

MT–52 Nuclear Scan Synthetic Origination

0/4

DOCUMENTED. S181b directive. mt_originator.py exists but needs extension.

MT-53 **Pokemon Autonomous Bot** 1/9

Phase 1 COMPLETE (S199). Research: research/MT53_POKEMON_RESEARCH.md. Verdict: PyBoy for GBC (Crystal + ROM hacks), mGBA + mGBA-http for GBA (Emerald). Rules-based engine (not LLM — Matthew: "doesn't burn tokens"). Key reference: llm_pokemon_scaffold (PyBoy + Claude function calling). Phase 2 next: emulator_control.py + state_reader.py + game_state.py. ROM hacks list TBD from Matthew.

Pending & Blocked

6 tasks awaiting start or unblocked

MT-5 **Claude Pro ↔ Claude Code Bridge** Not Started

Bridges Claude Pro web chat with Claude Code CLI sessions. Share context and continue conversations across interfaces.

MT-8 **iPhone Remote Control Perfection** Not Started

Control Claude Code sessions from your iPhone. Approve, reject, or steer work while away from your desk.

MT-18 **Academic Writing Workspace** Not Started

Academic writing workspace — structured paper composition with citation management. Professional research output.

MT-19 **Local LLM Fine-Tuning Capability** Not Started

Fine-tune local LLMs on CCA data. Long-term goal: specialized models trained on your specific workflows.

MT-25 **Presentation Generator** Blocked

Generate presentation slides matching Matthew's preferred style. One command from data to polished deck.

MT-36 **Session Efficiency Optimizer — Quality-Preserving Speed** Not Started

Measures and optimizes session overhead (init, wrap, tests) without cutting quality. Every second saved on overhead is a second spent shipping code.

PRIORITY QUEUE

Priority Queue

Active tasks ranked by decay-based priority scoring. Score = base value + aging penalty. Higher = work on this first.

1	MT-21	**Hivemind coordination** Phase 2 PASSED (6th). Phase 3 = 3-chat	10
2	MT-33	**Strategic Intelligence Report** Phase 4-6: chart integration, self-reference, hardening. kalshi_data_collector + learning_data_collector + report_differ built (S122-S123).	10
3	MT-23	**Mobile Remote Control v2** Research: evaluate Telegram/Discord channels MCP	9
4	MT-26	**Financial Intelligence Engine** CCA scope COMPLETE. Tier 3 Phase 2 deferred (needs numpy). 79 pipeline tests.	8
5	MT-32	**Visual Excellence & Design Engineering** Phases 1-3 COMPLETE. design_tokens canonical module, all consumers wired. Next: Phase 4 (UI components).	8
—	MT-34	**Medical AI (OpenEvidence replacement)** BLOCKED: Matthew refining concept. Do not start.	6
8	MT-29	**Cowork + Pro Bridge Hivemind** Research: evaluate Cowork, bridge Pro↔Code	5
9	MT-18	**Academic writing** Research: install/evaluate ClaudePrism	7
10	MT-13	**iOS/macOS app dev** Phase 3: first real app	7

LIVE INFRASTRUCTURE

Live Infrastructure

18 hooks integrated into Claude Code via settings.local.json. Every tool call passes through this pipeline.

Event	Matcher	Hook	Purpose
PostToolUse	*	meter.py	Token counter
PostToolUse	*	compact_anchor.py	Context anchor writes
PostToolUse	*	senior_dev_hook.py	SATD + effort + quality scoring
PostToolUse	*	queue_hook.py	Cross-chat queue context
PreToolUse	*	alert.py	Warn/block at red/critical
PreToolUse	*	cost_alert.py	Cost threshold warning
PreToolUse	*	path_validator.py	Dangerous path detection
PreToolUse	*	edit_guard.py	Edit retry prevention
PreToolUse	*	validate.py	Spec guard + plan compliance
PreToolUse	Bash	credential_guard.py	Credential extraction guard
PreToolUse	Bash	bash_guard.py	Command safety (network, packages, system)
UserPromptSubmit	*	skill_activator.py	Spec auto-activation
UserPromptSubmit	*	skillbook_inject.py	Strategy injection
UserPromptSubmit	*	capture_hook.py	Real-time memory capture
UserPromptSubmit	*	queue_hook.py	Cross-chat queue injection
Stop	*	auto_handoff.py	Block exit at critical
Stop	*	capture_hook.py	Session-end memory capture
PostCompact	*	post_compact.py	Recovery + journal logging

INTELLIGENCE & RESEARCH

Intelligence & Research

Reddit Intelligence

FINDINGS BY VERDICT

BUILD	33
ADAPT	71
REFERENCE	295
PERSONAL	0
SKIP	62

7 subreddits · 30 repos evaluated

Self-Learning System

418 PRINCIPLES

257/137 WINS / PAINS

Avg Score	0.48 (higher = more validated)
Journal Sessions	144 tracked
Papers Logged	**1242**
Sentinel Rate	0%

RESEARCH ROI **82** deliveries → **0** implemented → **0** profitable

Intelligence Verdicts

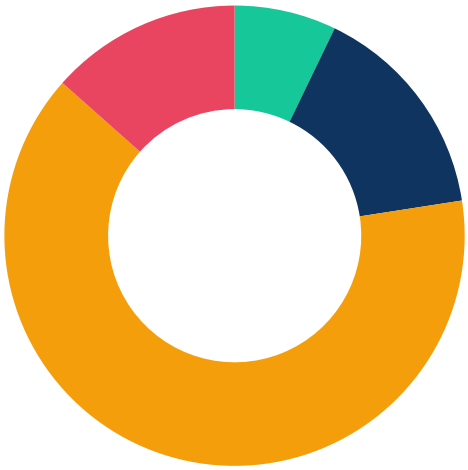

BUILD (33) ADAPT (71) REFERENCE (295)

FINANCIAL ANALYTICS — KALSHI BOT

Financial Analytics — Kalshi Bot

LIVE TRADES	WIN RATE	TOTAL P&L	SETTLED
1603	83.8%	$69.89	1588

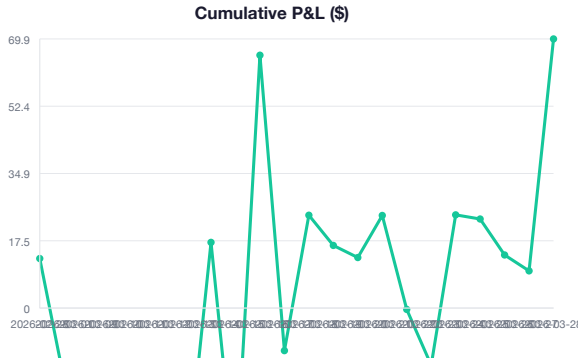

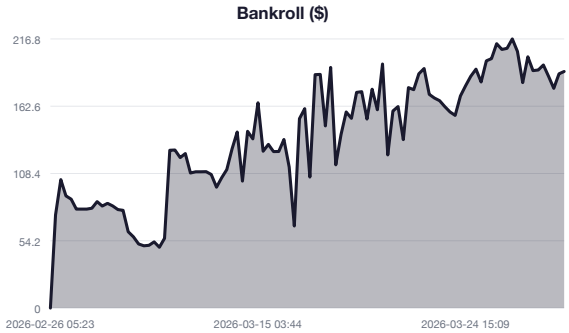

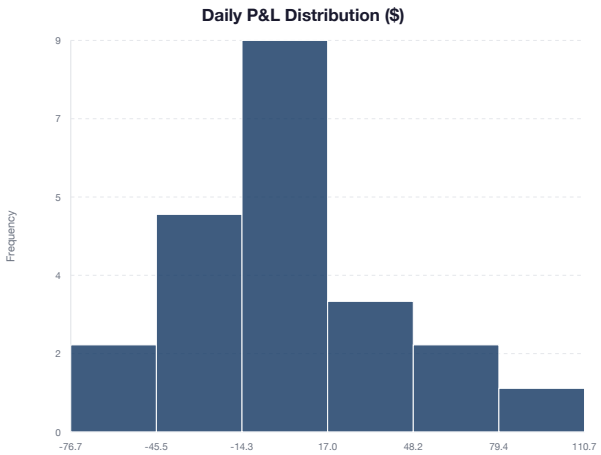
Daily P&L Distribution ($)
Frequency
9
7
5
4
2
0
-76.7
-45.5
-14.3
17.0
48.2
79.4
110.7

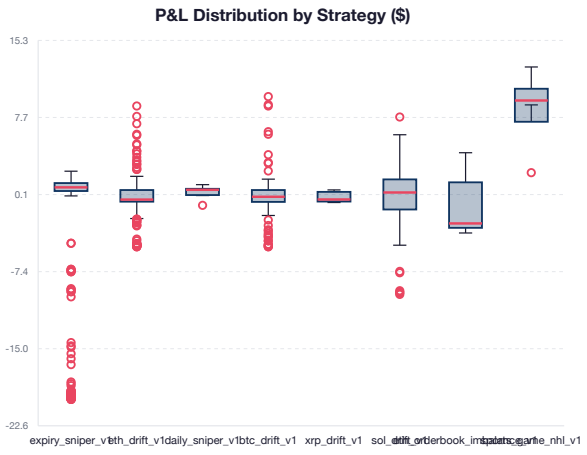
P&L Distribution by Strategy ($)
15.3
7.7
0.1
-7.4
-15.0
-22.6
expiry_sniper_v1
eth_drift_v1
daily_sniper_v1
btc_drift_v1
xrp_drift_v1
nhl_v1

Strategy Performance

Strategy	Trades	Wins	Win %	Avg P&L	Total P&L
expiry_sniper_v1	1125	1050	93.3%	$0.07	**$81.34**
daily_sniper_v1	116	100	99%	$0.51	**$51.58**
sports_game_nhl_v1	4	4	100%	$8.41	**$33.66**
btc_lag	1	1	100%	$3.5	**$3.5**
btc_lag_v1	1	1	100%	$0.57	**$0.57**
xrp_drift_v1	51	24	47.1%	$−0.07	**$−3.8**
eth_lag_v1	3	1	33.3%	$−2.18	**$−6.53**
btc_drift_v1	80	40	50%	$−0.12	**$−9.53**

sol_drift_v1	47	31	66%	$−0.33	**$−15.68**
eth_orderbook_imbalance_v1	15	5	33.3%	$−1.21	**$−18.2**
sports_game_nba_v1	2	0	0%	$−9.84	**$−19.67**
eth_drift_v1	158	73	46.2%	$−0.17	**$−27.35**

SELF-LEARNING INTELLIGENCE

Self-Learning Intelligence

JOURNAL ENTRIES	WINS / PAINS	CURRENT APF	POSTS REVIEWED
1232	257 / 137	22.9%	398

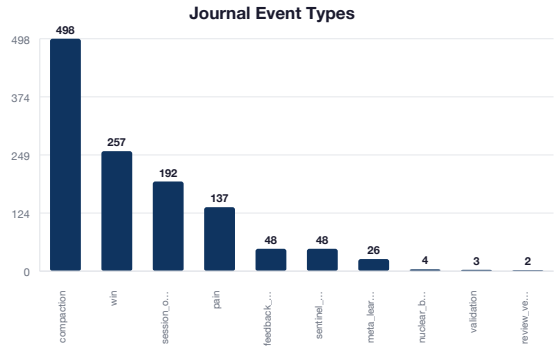

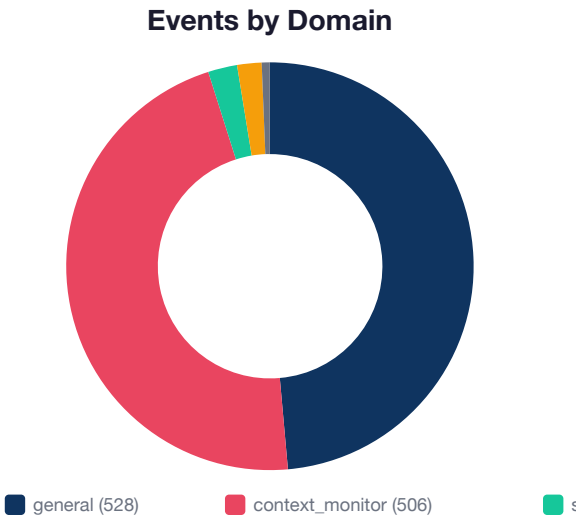

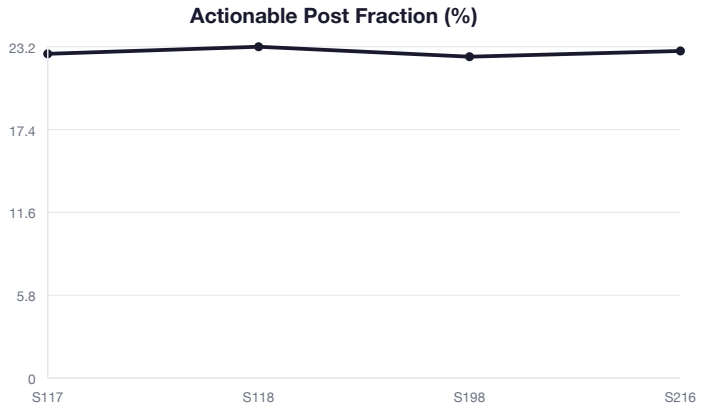

ARCHITECTURE DECISIONS

Architecture Decisions

Key technical decisions and their rationale. Each decision was made to optimize for simplicity, reliability, or developer experience.

Local-first storage

User owns all data, no cloud dependency

Stdlib-first (no pip)

Zero dependency management overhead

Hook-based delivery

Native Claude Code integration points

Adaptive context thresholds

Fixed percentages break at 1M windows

JSONL for all logging

Append-only, grep-friendly, no corruption risk

QualityGate geometric mean

Prevents Goodhart's Law gaming of metrics

Whitelist-first phishing detection

Zero false positives on legitimate domains

Typst over WeasyPrint

Single binary, millisecond compile, JSON-native

RISKS & BLOCKERS

Risks & Blockers

MT-25 Presentation Generator (Matthew's Style)

BLOCKER

NOT STARTED. BLOCKED on style samples from Matthew.

2 MTs stagnating

RISK

2 active tasks have negative stagnation scores — no progress in recent sessions.

Mitigation: Review priority picker output. Consider deprioritizing or archiving stalled work.

5 severity-3 learnings active

DEBT

High-severity patterns requiring ongoing vigilance across all sessions.

Mitigation: Encoded in global rules. Auto-enforced by hooks where possible.

NEXT PRIORITIES

Next Priorities

1. **Chat 15**
   cca-scout agent + test-runner hardening + Ebbinghaus decay
2. **Chat 16**
   Agent Teams in /cca-nuclear + SessionStart hook + SubagentStart budget
3. **Chat 17**
   Compaction v2 + cross-chat delivery + Phase 5 plan

HONEST ASSESSMENT

Honest Assessment

Objective gaps, limitations, and areas where the project falls short of its stated goals. Included for accountability.

18 master tasks stalled at Phase 1

GAP

MT-0, MT-9, MT-11, MT-12, MT-13, MT-27, MT-28, MT-30, MT-38, MT-39, MT-44, MT-46, MT-48, MT-49, MT-50, MT-51, MT-52, MT-53 — started but no subsequent phases shipped. Risk of scope sprawl without depth.

ACTION: Pick one stalled MT and advance it to Phase 2 this week. Suggested: MT-0 (oldest).

Single-developer project with no external users

LIMITATION

All tooling built for one developer. No community testing, no feedback loop from other Claude Code users. Limits validation of 'community-demanded' features.

ACTION: Consider publishing 1-2 standalone tools (e.g. context monitor, memory system) as open-source to get external feedback.

High test count (7,029) may overstate coverage

NUANCE

Average 42 tests per source file. Many test trivial properties (dict keys, return types). Integration test coverage across modules is limited to 1 suite.

ACTION: Add 3-5 cross-module integration tests that exercise real workflows (e.g. memory capture → retrieval, context alert → handoff).

1 task(s) blocked with no resolution timeline

BLOCKER

MT-25: Presentation Generator

ACTION: Review each blocked MT: either unblock with a workaround, rescope to remove the blocker, or archive if no longer relevant.

CI/CD pipeline exists but unverified

NUANCE

GitHub Actions workflow defined but no evidence of runs in git history. All testing is local-only.

ACTION: Push to GitHub and verify the Actions workflow runs. Fix any failures. Add status badge to README.

STATUS

All Systems Operational

10,706	1,902,376	246	1381
TESTS	LOC	SESSIONS	COMMITS

9 modules · 51 master tasks · 282 test suites · Zero external dependencies

10	35	461
MTs Complete	MTs Active	Findings

Generated 2026-04-01 · Session 246 · github.com/mpshields96/ClaudeCodeAdvancements